

AWS Auto Scaling Group with Application Loa

Introduction

Multi-Service Auto Scaling Configuration

In this project, different scaling methods are used for different services to demonstrate how AWS Auto Scaling works in real-world cloud environments.

Each service follows a separate scaling strategy based on application requirements and traffic patterns.

Service	Scaling Type	Description
Home Service	Dynamic Scaling	Automatically adds or removes EC2 instances depending on traffic and CPU usage.
Mobile Service	Dynamic + Scheduled Scaling	Uses automatic scaling during high traffic and scheduled scaling for fixed peak hours.
Laptop Service	Static Scaling	Runs with a fixed number of EC2 instances without automatic scaling.

Scaling Techniques Used

1. Dynamic Scaling — Home Service

The Home Service uses Dynamic Scaling to automatically manage EC2 instances according to application load.

How it works:

- When traffic or CPU usage increases, AWS automatically launches new EC2 instances.
- When traffic becomes low, extra EC2 instances are automatically removed.

Advantages:

- Better performance during high traffic
- Efficient resource utilization
- Reduced infrastructure cost
- Automatic scaling without manual intervention

2. Dynamic + Scheduled Scaling — Mobile Service

The Mobile Service uses both Dynamic Scaling and Scheduled Scaling to handle variable and predictable traffic.

Dynamic Scaling

AWS automatically increases or decreases EC2 instances based on CPU usage and incoming traffic.

Scheduled Scaling

AWS launches or terminates instances at predefined times according to expected traffic patterns.

Example:

- During busy hours, more EC2 instances are launched automatically.
- During low traffic periods, unnecessary instances are removed.

Advantages:

- Handles expected and unexpected traffic efficiently
- Better application availability
- Improved performance during peak hours
- Optimized cloud resource usage

3. Static Scaling — Laptop Service

The Laptop Service uses Static Scaling with a fixed number of EC2 instances.


How it works:


- A predefined number of EC2 instances always remain running.
- No automatic scaling operations are performed.

Advantages:

- Stable and predictable infrastructure
- Simple configuration and management
- Suitable for applications with constant traffic

Extended Architecture







AWS Services Used

Service	Purpose
EC2	Hosting Applications
Auto Scaling Group	Automatic Scaling
Launch Template	EC2 Configuration
Application Load Balancer	Traffic Distribution
Target Groups	Route Requests
CloudWatch	Monitoring
Security Groups	Allow HTTP Traffic

Extended Project Structure





Step 1 — Create Launch Templates

Create separate Launch Templates for:

- Home Service
- Mobile Service
- Laptop Service

Configuration:

- Amazon Linux 2023
- t3.micro
- Security Group
- Key Pair
- User Data Script

Step 2 — Create Target Groups

Create three target groups:

Target Group	Service
home-tg	Home Service
mobile-tg	Mobile Service
laptop-tg	Laptop Service

Configuration:

- Protocol → HTTP
- Port → 80
- Target Type → Instance

Step 3 — Create Application Load Balancer

Create an Application Load Balancer with:

- Internet Facing
- HTTP Listener on Port 80
- Attach all Target Groups

Step 4 — Configure Listener Rules

Add path-based routing rules:

Path	Target Group
/mobile/*	mobile-tg
/laptop/*	laptop-tg
Default	home-tg

Step 5 — Create Auto Scaling Groups

Create three Auto Scaling Groups:

Auto Scaling Group	Desired Capacity
home-asg	2
mobile-asg	2
laptop-asg	2

Attach corresponding:

- Launch Template
- Target Group
- Load Balancer

Auto Scaling Working

The Auto Scaling Group continuously monitors EC2 instance health and automatically maintains the required number of instances.

If one instance fails:

- New EC2 instance is automatically launched.

If traffic increases:

- Additional EC2 instances can be created automatically.

Benefits:

- High Availability
- Fault Tolerance
- Better Performance
- Automatic Recovery
- Scalability

Round Robin Technique

The Application Load Balancer uses the Round Robin algorithm to distribute incoming traffic equally between multiple EC2 instances.

Example:

- Request 1 → EC2-1
- Request 2 → EC2-2
- Request 3 → EC2-1
- Request 4 → EC2-2

This helps in:

- Balanced Traffic Distribution
- Reduced Server Load
- Improved Application Performance
- Better Resource Utilization

Application URLs

Home Page

```
http://application-lb-dns-name/
```

Mobile Page


```
http://application-lb-dns-name/mobile/
```

Laptop Page


```
http://application-lb-dns-name/laptop/
```

Screenshots


Auto Scaling Groups



EC2 Running Instances




Home Page Output




This is Home page ip-172-31-95-238.ec2.internal

Mobile Page Output



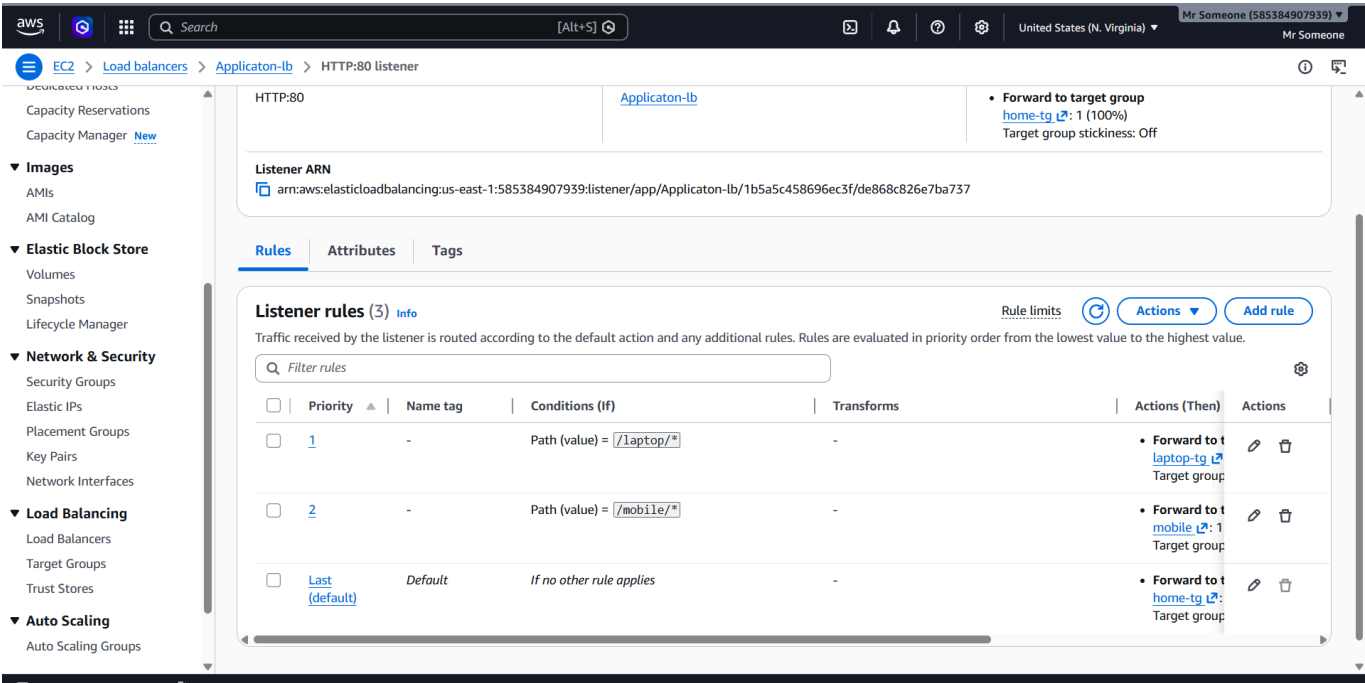
This is mobile page ip-172-31-91-31.ec2.internal

Laptop Page Output




This is laptop page ip-172-31-6-130.ec2.internal

Load Balancer Listener Rules



Target Groups



Security Group Configuration

Type	Port
SSH	22
HTTP	80

Features

- Auto Scaling Groups
- Multiple EC2 Instances
- Dynamic Scaling
- Application Load Balancer
- Path-Based Routing
- Round Robin Load Distribution
- High Availability
- Fault Tolerance
- Automatic Recovery
- Scalable Cloud Architecture

Future Improvements

- HTTPS SSL Configuration
- Route53 Domain Setup
- Jenkins CI/CD Integration
- Docker Deployment
- Kubernetes Integration
- Multi-AZ Deployment
- CloudWatch Dashboard

Author

Aryan Rajendra Dhawas

Fortune Cloude Technology
AWS & DevOps Engineering

Summary

This project successfully demonstrates implementation of AWS Auto Scaling Groups integrated with an Application Load Balancer using extended cloud architecture. Separate Auto Scaling Groups were created for Home, Mobile, and Laptop services, with multiple EC2 instances attached to each service for scalability and high availability. The Application Load Balancer distributed traffic using path-based routing and the Round Robin technique, while Auto Scaling Groups automatically maintained healthy EC2 instances and supported dynamic scaling based on infrastructure requirements.
